


CLM-ml-jax GPU throughput: float32 vs float64 — CHATS7 · Quadro RTX 8000
(median of 10 vmapped forward passes per N , after JIT warm-up)

(a) GPU throughput — float32 vs float64



(b) XLA compile time vs N

